


Guard, explained in one picture

An agent proposes. Guard decides. A checkpoint enforces. A receipt explains what happened.




A receipt is evidence

Linked receipts help reveal changes to a task's history. Sensitive content is redacted before storage.

Five possible decisions

When rules disagree, the stricter outcome wins. Strictest is on the left.



A receipt accompanies every outcome in this design. This is a logical architecture, not a network deployment map.

The story: an assistant asks for too much

A hospital assistant tries to read every patient record. Follow the request from intent to evidence.

1 The agent proposes a query

It asks the database tool to read every patient record. This is the request Guard must evaluate.

```
SELECT * FROM patients
```

2 The checkpoint catches it

Before the tool runs, its checkpoint gathers the agent identity, workspace, tool name, and exact query.

```
WHO + WHAT + WHERE
```

3 Guard checks the rulebook

The installed hospital pack contains a rule against this broad query. The shared engine returns Block.

```
DECISION: BLOCK
```

4 The checkpoint stops it

The query is never sent to the database. The checkpoint returns the reason to the agent.

```
NO DATABASE EXECUTION
```

5 The team is notified

In the document's example, the decision also requires the checkpoint to notify the compliance team.

```
A REQUIRED FOLLOW-UP
```

6 A receipt tells the story

It records the blocked query attempt, the rule, the time, and whether notification completed.

```
LINKED TO THIS TASK'S HISTORY
```

Say it aloud: "The agent asked. Guard said no. The checkpoint stopped it and kept the reason."

Three kinds of checks use the same architecture

ACTION

May the agent do this?

Before a tool operates.

Tool gateway (MCP), workflow runtime, and CLI hook.

PROMPT

May this go to the model?

Before a prompt reaches a model.

LLM proxy and the LiteLLM plugin's pre-call check.

RESPONSE

May this return to the caller?

Before model output is delivered.

LLM proxy. LiteLLM response support is deferred.

Each request uses the applicable gate. Lens uses the existing tool and proxy paths.

Response protection depends on the integration and streaming mode. Phase 5 task permissions are shown separately on page 1.